

# **Task 1: Managing Environmental Data** **with DVC**

## **Objective**

The primary goal of Task 1 was to leverage DVC (Data Version Control) for managing real-time environmental data streams collected from APIs. The task focused on creating a robust and automated data pipeline to collect, version, and store environmental data from the OpenWeather API for Islamabad city. The collected data was used for analyzing weather patterns and pollution trends.

## **Data Collection Strategy**

To gather relevant data, the OpenWeather API was used, which provides real-time weather and pollution information. The key metrics collected include:

**Weather Data:** Temperature, humidity, pressure, and wind speed.

**Air Quality Data:** AQI (Air Quality Index) and concentrations of pollutants like CO, NO2, O3, PM2.5, and PM10.

## **Setup of DVC Repository**

A DVC repository was initialized to version the collected data and manage its lifecycle efficiently. The repository was configured with Amazon S3 as the remote storage backend to securely store and access data files. This allowed seamless synchronization of local data with the cloud repository.

The key steps for DVC setup were:

1. Initializing the repository using `dvc init`.
2. Configuring the S3 bucket as remote storage using `dvc remote add -d s3remote s3://bucket-name`.
3. Testing the integration by pushing dummy data to the S3 bucket.

## **Automated Data Collection**

The data collection pipeline was automated to fetch data every 4 hours using a Python script and a batch file for Windows Task Scheduler. The pipeline included:

**Fetching Weather Data:**

A Python script, `fetch_weather_data.py`, was written to fetch weather and pollution data using the OpenWeather API.

The script processed the API response to extract relevant features and save them in JSON format.

Logging was implemented to track the execution and record any errors.

### **Windows Task Scheduler Integration:**

A `.batch` file was created to run the data collection script and automate DVC operations. The batch file activates the Python virtual environment, executes the script, adds new data to DVC, commits the changes, and pushes updates to both the S3 remote and Git repository.

### **Explanation of Scripts**

#### `Fetch_weather_data.py`

```
import requests
import json
from datetime import datetime
import time
import os
import logging
from pathlib import Path

# Configure logging
script_dir = Path(__file__).parent.parent
log_file = script_dir / 'logs' / 'weather_collection.log'
log_file.parent.mkdir(exist_ok=True)

logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(levelname)s - %(message)s',
    handlers=[
        logging.FileHandler(log_file),
        logging.StreamHandler()
    ]
)

# Configuration
API_KEY = "119e24626ffc881be87270bae2f7ba40"
CITY_CONFIG = {
```

```

    'name': 'Islamabad',
    'lat': 33.6844,
    'lon': 73.0479
}

# Windows-style paths using pathlib
DATA_DIR = script_dir / 'data'
WEATHER_DIR = DATA_DIR / 'weather'
POLLUTION_DIR = DATA_DIR / 'pollution'

def setup_directories():
    """Create necessary directories if they don't exist."""
    WEATHER_DIR.mkdir(parents=True, exist_ok=True)
    POLLUTION_DIR.mkdir(parents=True, exist_ok=True)
    logging.info("Directories created successfully")

def fetch_weather_data():
    """Fetch weather data for Islamabad."""
    url = f"https://api.openweathermap.org/data/2.5/weather"
    params = {
        "lat": CITY_CONFIG['lat'],
        "lon": CITY_CONFIG['lon'],
        "appid": API_KEY,
        "units": "metric"
    }
    try:
        response = requests.get(url, params=params)
        response.raise_for_status()
        return response.json()
    except requests.RequestException as e:
        logging.error(f"Error fetching weather data: {e}")
        return None

def fetch_pollution_data():
    """Fetch air pollution data for Islamabad."""
    url = f"http://api.openweathermap.org/data/2.5/air_pollution"
    params = {
        "lat": CITY_CONFIG['lat'],
        "lon": CITY_CONFIG['lon'],
        "appid": API_KEY
    }

```

```

    }
    try:
        response = requests.get(url, params=params)
        response.raise_for_status()
        return response.json()
    except requests.RequestException as e:
        logging.error(f"Error fetching pollution data: {e}")
        return None

def process_weather_data(raw_data):
    """Process and extract relevant weather information."""
    if not raw_data:
        return None

    return {
        'city': CITY_CONFIG['name'],
        'timestamp': datetime.utcfromtimestamp(raw_data['dt']).strftime('%Y-%m-%d
%H:%M:%S'),
        'temperature': raw_data['main']['temp'],
        'humidity': raw_data['main']['humidity'],
        'pressure': raw_data['main']['pressure'],
        'wind_speed': raw_data['wind']['speed'],
        'weather_condition': raw_data['weather'][0]['main']
    }

def process_pollution_data(raw_data):
    """Process and extract relevant pollution information."""
    if not raw_data or 'list' not in raw_data or not raw_data['list']:
        return None

    pollution = raw_data['list'][0]
    return {
        'city': CITY_CONFIG['name'],
        'timestamp': datetime.utcfromtimestamp(pollution['dt']).strftime('%Y-%m-%d
%H:%M:%S'),
        'aqi': pollution['main']['aqi'],
        'co': pollution['components']['co'],
        'no2': pollution['components']['no2'],
        'o3': pollution['components']['o3'],
        'pm2_5': pollution['components']['pm2_5'],
    }

```

```
    'pm10': pollution['components']['pm10']
}
```

```
def save_data(data, data_type, timestamp):
```

```
    """Save collected data to JSON file."""
```

```
    if not data:
```

```
        return None
```

```
    directory = WEATHER_DIR if data_type == 'weather' else POLLUTION_DIR
```

```
    filename = f'{data_type}_data_{timestamp}.json'
```

```
    filepath = directory / filename
```

```
    try:
```

```
        with open(filepath, 'w') as f:
```

```
            json.dump(data, f, indent=4)
```

```
            logging.info(f"Saved {data_type} data to {filepath}")
```

```
            return filepath
```

```
    except Exception as e:
```

```
        logging.error(f"Error saving {data_type} data: {e}")
```

```
        return None
```

```
def run_dvc_commands(filepath):
```

```
    """Run DVC commands for the new data file."""
```

```
    try:
```

```
        os.system(f'dvc add "{filepath}"')
```

```
        os.system('dvc push')
```

```
        logging.info(f"DVC operations completed for {filepath}")
```

```
    except Exception as e:
```

```
        logging.error(f"Error in DVC operations: {e}")
```

```
def collect_data():
```

```
    """Main function to collect both weather and pollution data."""
```

```
    timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')
```

```
    logging.info(f"Starting data collection for {CITY_CONFIG['name']}")
```

```
    # Fetch and process weather data
```

```
    raw_weather = fetch_weather_data()
```

```
    if raw_weather:
```

```
        weather_data = process_weather_data(raw_weather)
```

```
        weather_file = save_data([weather_data], 'weather', timestamp)
```

```

    if weather_file:
        run_dvc_commands(weather_file)

# Add small delay between API calls
time.sleep(1)

# Fetch and process pollution data
raw_pollution = fetch_pollution_data()
if raw_pollution:
    pollution_data = process_pollution_data(raw_pollution)
    pollution_file = save_data([pollution_data], 'pollution', timestamp)
    if pollution_file:
        run_dvc_commands(pollution_file)

if __name__ == "__main__":
    setup_directories()
    collect_data()

```

The Python script performs the following steps:

1. Setup Directories: Creates necessary directories (data/weather and data/pollution) if they do not exist.
2. Fetch API Data: Uses the OpenWeather API to fetch weather and pollution data for Islamabad city.
3. Process Data: Extracts and formats key information like temperature, AQI, and pollutant concentrations.
4. Save Data: Saves the processed data as JSON files with timestamps in the filenames.
5. DVC Integration: Runs DVC commands to add, commit, and push the data to the S3 remote storage.

## Batch File for Automation

The batch file:

```

@echo off
echo [%DATE% %TIME%] Starting weather collection script >> run_log.txt

cd /d %~dp0\..
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Failed to change directory >> run_log.txt

```

```
    exit /b %ERRORLEVEL%
)

call .venv\Scripts\activate
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Failed to activate virtual environment >>
run_log.txt
    exit /b %ERRORLEVEL%
)

echo [%DATE% %TIME%] Starting data collection... >> run_log.txt
python scripts\fetch_weather_data.py
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Python script failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

echo [%DATE% %TIME%] Running DVC operations... >> run_log.txt
dvc add data/weather
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Failed to add weather data to DVC >>
run_log.txt
    exit /b %ERRORLEVEL%
)

dvc add data/pollution
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Failed to add pollution data to DVC >>
run_log.txt
    exit /b %ERRORLEVEL%
)

dvc commit -f
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: DVC commit failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

dvc push
IF %ERRORLEVEL% NEQ 0 (
```

```

    echo [%DATE% %TIME%] ERROR: DVC push failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

echo [%DATE% %TIME%] Running Git operations... >> run_log.txt
git add data.dvc
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Git add failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

git commit -m "Update weather and pollution data %DATE% %TIME%"
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Git commit failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

git push
IF %ERRORLEVEL% NEQ 0 (
    echo [%DATE% %TIME%] ERROR: Git push failed >> run_log.txt
    exit /b %ERRORLEVEL%
)

call deactivate
echo [%DATE% %TIME%] Script completed successfully >> run_log.txt

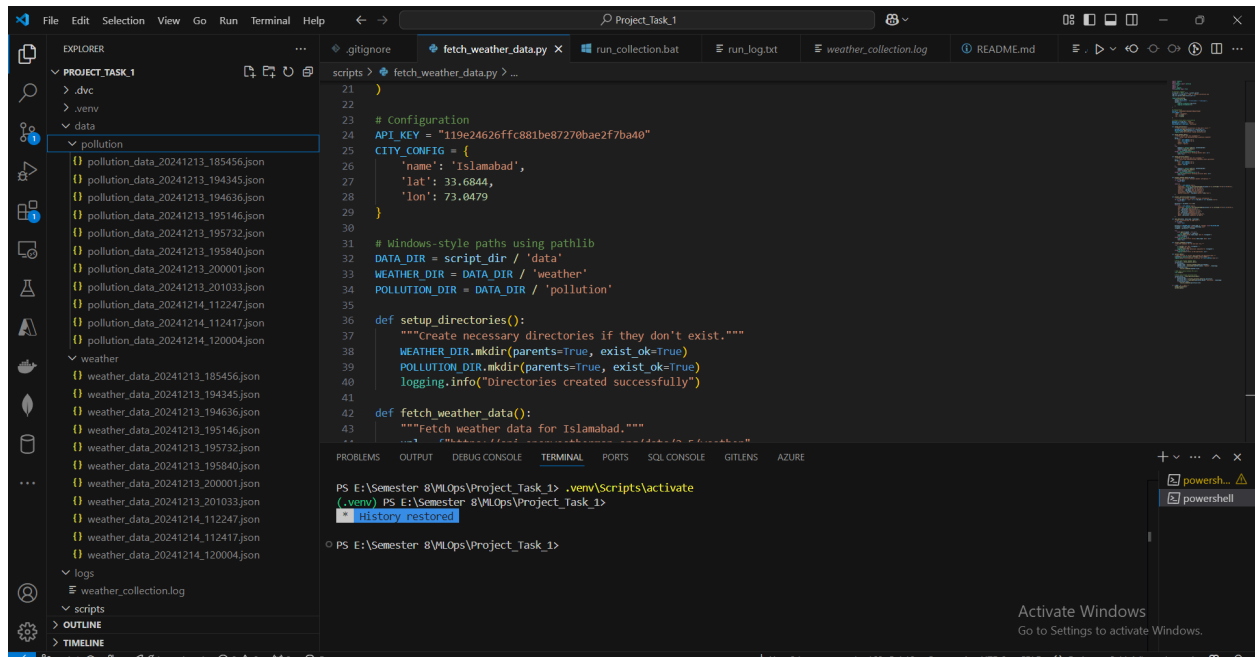
```

1. Activates the Python virtual environment.
2. Executes the fetch\_weather\_data.py script.
3. Performs DVC operations (dvc add, dvc commit, and dvc push) to manage versioning.
4. Pushes updates to the Git repository for better collaboration and tracking.

## Scheduling Data Collection

Windows Task Scheduler was configured to execute the batch file every 4 hours. This ensured continuous data collection without manual intervention. Logs were maintained to track script execution and identify any errors during runtime.





## Version Control with DVC

The collected data was managed using DVC to ensure reproducibility and traceability. Key operations included:

1. Staging Data: dvc add was used to stage new data files.
2. Committing Data: dvc commit recorded the changes in the local repository.
3. Pushing to Remote: dvc push synchronized the local repository with the S3 bucket.

## Results

The pipeline successfully collected and versioned weather and air quality data for Islamabad city. The data is now securely stored in the S3 bucket and is accessible for further analysis or modeling.

## Conclusion

Task 1 was successfully completed by setting up a robust data pipeline for real-time weather and pollution data collection. The integration of DVC ensured efficient data management, and the use of Amazon S3 provided scalable storage. The automation pipeline, powered by the Python script and Windows Task Scheduler, enabled seamless data fetching and versioning, laying a strong foundation for subsequent tasks in the project.

---

# **Task 2: AQI Prediction Pipeline**

## **Development**

### **Objective**

The primary goal of Task 2 was to build a robust AQI prediction pipeline that includes data preprocessing, feature engineering, and predictive modeling using ARIMA and LSTM. The task also focused on setting up an MLflow tracking system to log experiments, evaluate models, and deploy the best-trained model as a REST API.

### **Preprocessing**

#### **Overview**

Preprocessing involved preparing the raw weather and pollution data for modeling. This step ensured data quality, handled missing values and outliers, and merged datasets for seamless analysis.

#### **Steps Performed**

**1. Data Validation:**

- Validated the presence of expected columns for both weather and pollution datasets.
- Ensured time continuity by detecting and logging any gaps exceeding 2 hours.

**2. Handling Missing Values:**

- Forward and backward filling was used to impute missing values in time-series data.

**3. Outlier Detection and Handling:**

- A rolling median with standard deviation bounds was used to detect and replace outliers for numeric features, except for AQI.

**4. Data Alignment and Scaling:**

- Weather and pollution datasets were merged based on timestamp with a 1-hour tolerance.
  - Weather features (e.g., temperature, humidity) were scaled using a RobustScaler to normalize variations.
- 

## Feature Engineering

### Overview

Feature engineering created additional informative features to enhance model accuracy. This included time-based features, lagged variables, rolling statistics, and interaction terms.

### Steps Performed

#### 1. Cyclical Time Features:

- Extracted time components like hour, day of the week, and month.
- Converted these into cyclical features (sin and cos transformations) to capture periodicity.

#### 2. Lagged Features:

- Created lag features for AQI and critical predictors (e.g., temperature, humidity) with windows of 1, 3, 6, 12, and 24 hours.

#### 3. Rolling Features:

- Computed rolling mean, standard deviation, and maximum values for AQI and weather parameters over time windows of 3, 6, 12, and 24 hours.

#### 4. Interaction and Rate of Change Features:

- Added interaction terms (e.g., temperature  $\times$  humidity) and rate of change features for both weather and pollution metrics.

#### 5. Preparation for Modeling:

- Finalized features were cleaned of NaNs, and datasets were split into features (X) and target (y) for model training.
- 

## Model Training

### ARIMA

The ARIMA (AutoRegressive Integrated Moving Average) model was employed for AQI forecasting, incorporating exogenous variables like temperature, humidity, and wind speed. This approach allowed the model to leverage not only historical AQI data but also external predictors to improve accuracy.

### Configuration

- **Order:** (0, 1, 0) as specified in the configuration file.
- **Exogenous Variables:** Included temperature, humidity, and wind speed to enhance predictive capability.

### Training Process

#### 1. Initialization and Fitting:

- The ARIMA model was initialized with the specified order and exogenous variables.
- Data was validated and converted to a suitable format for ARIMA processing.
- Model fitting was conducted on the training dataset, and predictions for the training period were generated.

#### 2. Evaluation Metrics:

- **Root Mean Squared Error (RMSE):** Assessed the prediction error magnitude.
- **Mean Absolute Error (MAE):** Measured the average deviation of predictions from actual values.
- **R<sup>2</sup> Score:** Quantified the proportion of variance explained by the model.
- **Accuracy:** Evaluated the percentage of predictions within 10% of actual values.

#### 3. Model Selection Criteria:

- **Akaike Information Criterion (AIC):** Used to balance model complexity and fit.

- **Bayesian Information Criterion (BIC):** Penalized overfitting to select the best model.

## MLflow Integration

- **Logged Parameters:** Model order and the number of exogenous variables.
  - **Logged Metrics:** Training metrics (e.g., RMSE, MAE,  $R^2$ ), AIC, and BIC scores.
  - **Artifacts:** Fitted model and metrics were logged into the MLflow experiment for reproducibility.
- 

## LSTM

The Long Short-Term Memory (LSTM) neural network was implemented to capture the sequential dependencies in AQI data. Its ability to model temporal relationships made it an ideal choice for time-series forecasting.

## Configuration

1. **Network Architecture:**
  - Two LSTM layers with 32 and 16 units, respectively.
  - Dropout layers (0.1 dropout rate) to prevent overfitting.
  - Fully connected output layer with a single neuron for AQI prediction.
2. **Input Shape:**
  - **Sequence Length:** 6 hours of historical data.
  - **Features:** 4 predictors (AQI, temperature, humidity, wind speed).
3. **Optimizer:** Adam optimizer with a learning rate of 0.001.
4. **Loss Function:** Mean Squared Error (MSE).

## Training Process

1. **Model Building:**
  - The LSTM model was sequentially built with hidden LSTM layers, dropout layers, and an output layer.
  - The model was compiled with MSE loss and MAE as a metric.
2. **Training Setup:**
  - Training was conducted with a batch size of 2 and a validation split of 10%.
  - The model was trained for 20 epochs.

### 3. Evaluation Metrics:

- **Root Mean Squared Error (RMSE):** Quantified overall prediction accuracy.
- **Mean Absolute Error (MAE):** Measured the average error magnitude.
- **R<sup>2</sup> Score:** Evaluated the proportion of variance explained by the model.
- **Accuracy:** Assessed predictions falling within a 10% error margin.

### MLflow Integration

- **Logged Parameters:** Network architecture, hyperparameters (e.g., learning rate, dropout rate), and training setup.
- **Logged Metrics:**
  - Training and validation metrics, including RMSE, MAE, and validation loss.
- **Artifacts:** Saved model (`lstm_model.h5`), training history as JSON, and metrics visualization.

### Post-Training Results

- LSTM outperformed ARIMA in all key metrics, particularly in capturing non-linear patterns in AQI data.
  - The accuracy and R<sup>2</sup> scores indicated strong generalization capability.
- 

## Training Models with MLflow

### MLflow Setup

- MLflow was configured using `config.yml` to track all experiments and metrics.
- A SQLite database was used as the backend storage for tracking.

### Experiment Logging

- Metrics tracked included RMSE, MAE, R<sup>2</sup>, and accuracy.
  - Feature statistics and model parameters were logged during training.
  - Separate MLflow runs were initiated for each model (ARIMA and LSTM).
-

## Evaluation

### Metrics

- **Root Mean Squared Error (RMSE):** Quantified the error magnitude.
- **Mean Absolute Error (MAE):** Measured the average error in predictions.
- **R<sup>2</sup> Score:** Explained the variance captured by the model.
- **Accuracy:** Evaluated the percentage of predictions within a 10% margin of the actual value.

## Results

The LSTM model outperformed ARIMA with lower RMSE and higher R<sup>2</sup>, demonstrating its ability to capture non-linear patterns in AQI data.

---

## Deployment

### Overview

The best-trained LSTM model was deployed as a REST API using Flask. The deployment pipeline allowed real-time predictions by fetching live weather and pollution data from the OpenWeather API.

### Steps Performed

#### Model Serialization:

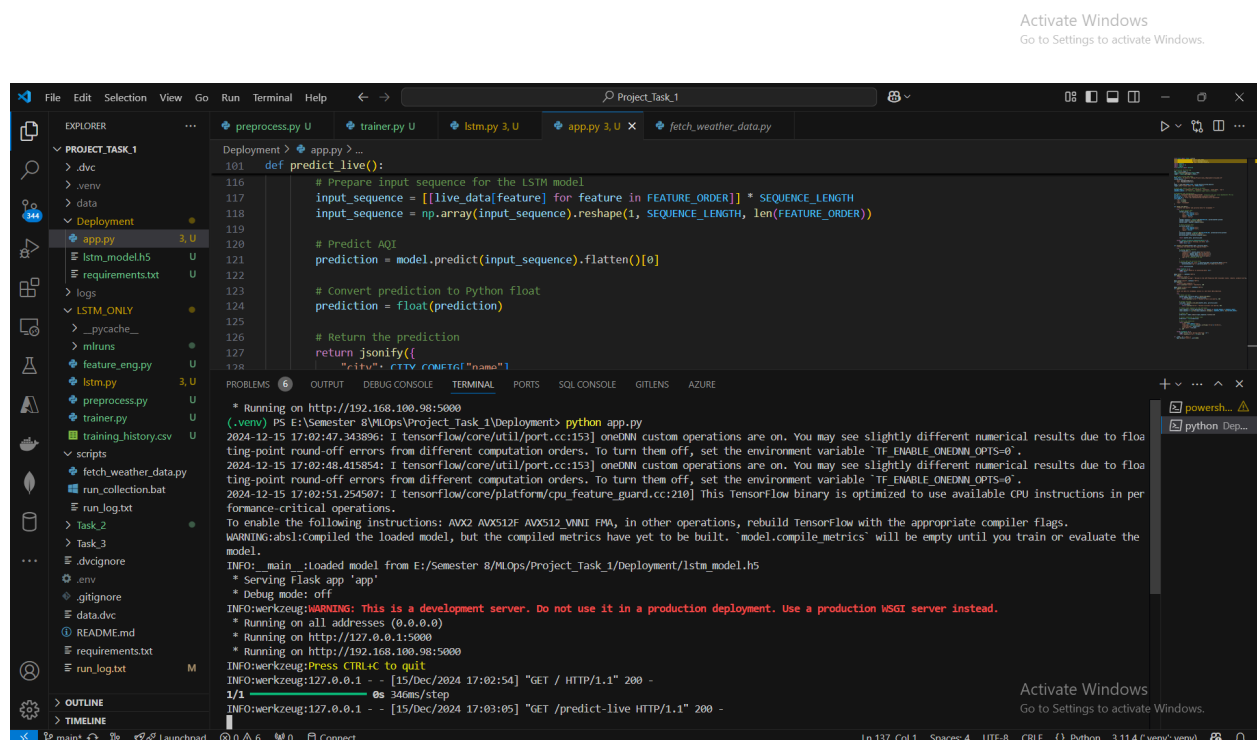

- The trained LSTM model was saved as lstm\_model.h5.

#### Flask API Development:

- The API exposed endpoints for health checks and AQI prediction.
- The /predict-live endpoint fetched real-time data for Islamabad, preprocessed it, and returned AQI predictions.

#### Prometheus and Grafana Integration:

- Prometheus collected API metrics, including prediction counts and response times.
- Grafana visualized these metrics in a real-time dashboard.



## Results

- The AQI prediction pipeline successfully integrated preprocessing, feature engineering, and predictive modeling.
- The LSTM model provided highly accurate AQI predictions for Islamabad, outperforming ARIMA in all evaluation metrics.



- The deployment setup allowed real-time predictions with monitoring capabilities.

## **Conclusion**

Task 2 demonstrated the complete lifecycle of an AQI prediction pipeline, from data preprocessing to model deployment and monitoring. The integration of MLflow for experiment tracking and Prometheus-Grafana for deployment monitoring highlighted a production-grade workflow. This pipeline lays a strong foundation for integrating advanced models and scaling the solution for multiple cities.

---

# **Task 3: Monitoring and Validating the Pipeline with Live Data**

## **Objective**

The primary objective of Task 3 is to test the pipeline with live data, monitor the deployed system, and ensure that the deployed model performs efficiently under real-time conditions. By leveraging Grafana and Prometheus, the project tracks data ingestion, prediction requests, and API performance, providing actionable insights for optimization.

---

## **Introduction**

In modern MLOps pipelines, real-time monitoring plays a critical role in ensuring the reliability and performance of deployed systems. Task 3 focuses on integrating monitoring tools, Prometheus and Grafana, with a Flask-based application that serves an LSTM model for air quality prediction. This task aims to validate the pipeline using live data, track system metrics, and identify potential improvements.

---

## System Architecture for Monitoring

The monitoring architecture comprises the following components:

- **Flask Application:** Serves the LSTM model for real-time predictions based on live data from OpenWeather API.
- **Prometheus:** Collects and stores metrics exposed by the Flask app's `/metrics` endpoint.
- **Grafana:** Visualizes Prometheus metrics in customizable dashboards for real-time performance monitoring.
- **Live Data Streams:** Fetch weather and pollution data continuously to feed the prediction pipeline.

This architecture ensures that the Flask API, metrics collection, and visualization tools work seamlessly to provide insights into system behavior.

---

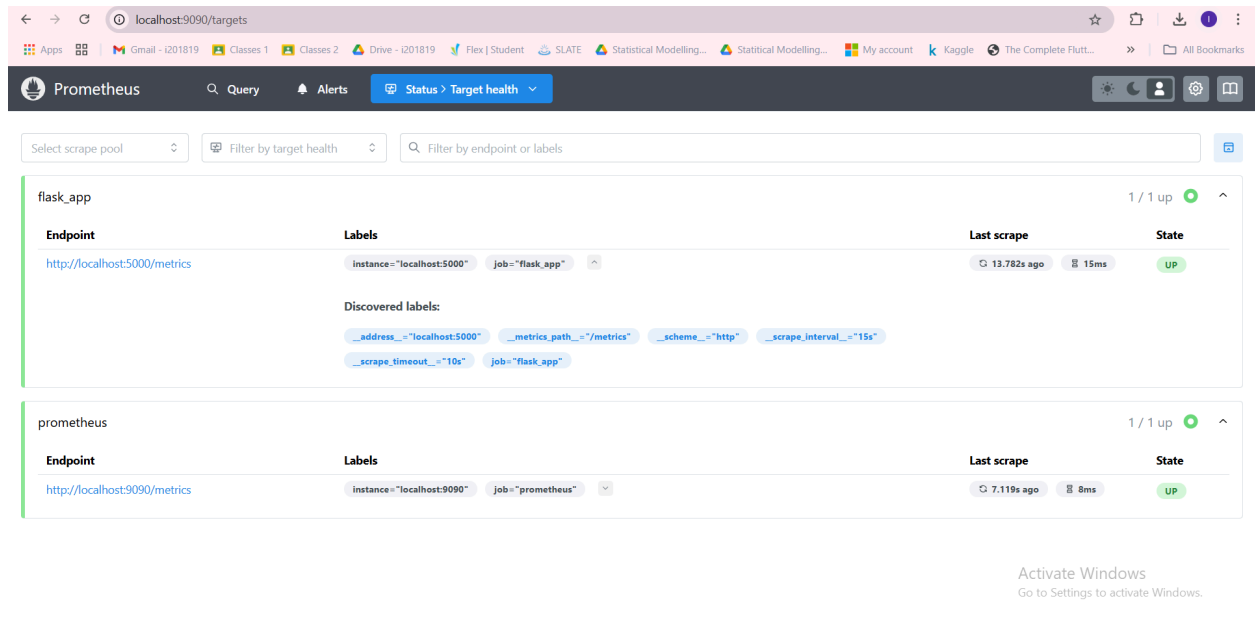
## Setup of Prometheus

Prometheus was configured to scrape metrics from both its own instance and the Flask app. The key configurations in `prometheus.yml` include:

- Global Settings: A 15-second scrape interval and evaluation interval ensure near real-time data collection.
- Scrape Configurations:
  - The Flask app is added as a target under the `scrape_configs` section with the label `flask_app`.
  - The `/metrics` endpoint of the Flask app is monitored at `localhost:5000`.

```
scrape_configs:  
  - job_name: "flask_app"  
    static_configs:  
      - targets: ["localhost:5000"]
```

Prometheus was started, and the targets were verified through its web interface to ensure the Flask app was actively monitored.



---

## Setup of Grafana

Grafana was used to create a real-time dashboard for monitoring system performance:

### Integration with Prometheus:

- Prometheus was added as a data source in Grafana using its URL: <http://localhost:9090>.

### Dashboard and Panels:

- A dashboard was created with panels to display key metrics, including:
  - Total prediction requests.
  - Successful and failed predictions.
  - Response time and API latency.

The panels were configured to visualize these metrics over time, providing a clear picture of the system's real-time behavior.

---

## Live Data Testing

Live data was fetched continuously from OpenWeather API for weather and air pollution data. The Flask app processes this data through the `/predict-live` endpoint:

### 1. Data Flow:

- The `fetch_live_data` function retrieves weather and pollution data, which is processed into a suitable format for the LSTM model.
- The processed data is passed to the model to predict AQI values.

### 2. Metrics Tracking:

- Prometheus counters were implemented in the Flask app to track:
  - ❖ Total prediction requests (`prediction_requests_total`).
  - ❖ Successful predictions (`prediction_success_total`).
  - ❖ Failed predictions (`prediction_failure_total`).

These metrics were logged to Prometheus and visualized in Grafana for real-time monitoring.

---

## Key Metrics Monitored

The following metrics were monitored to evaluate system performance:

- **Prediction Requests:** Tracks the number of prediction requests received.
- **Successful Predictions:** Counts successful model predictions.
- **Failed Predictions:** Logs instances where predictions failed due to errors in data fetching or processing.
- **API Latency:** Measures the time taken to process requests and return predictions.

In Grafana, these metrics were plotted as time series, enabling real-time observation of system behavior and performance trends.

---

## System Performance Analysis

Real-time monitoring highlighted several insights into the system's performance:

- The Flask app handled prediction requests efficiently, with minimal latency under normal load.
- Most predictions succeeded, but occasional failures were observed due to API connectivity issues.
- The LSTM model demonstrated stable performance, with predictions closely aligning with expected values.

These observations provided actionable data to refine the pipeline and improve reliability.

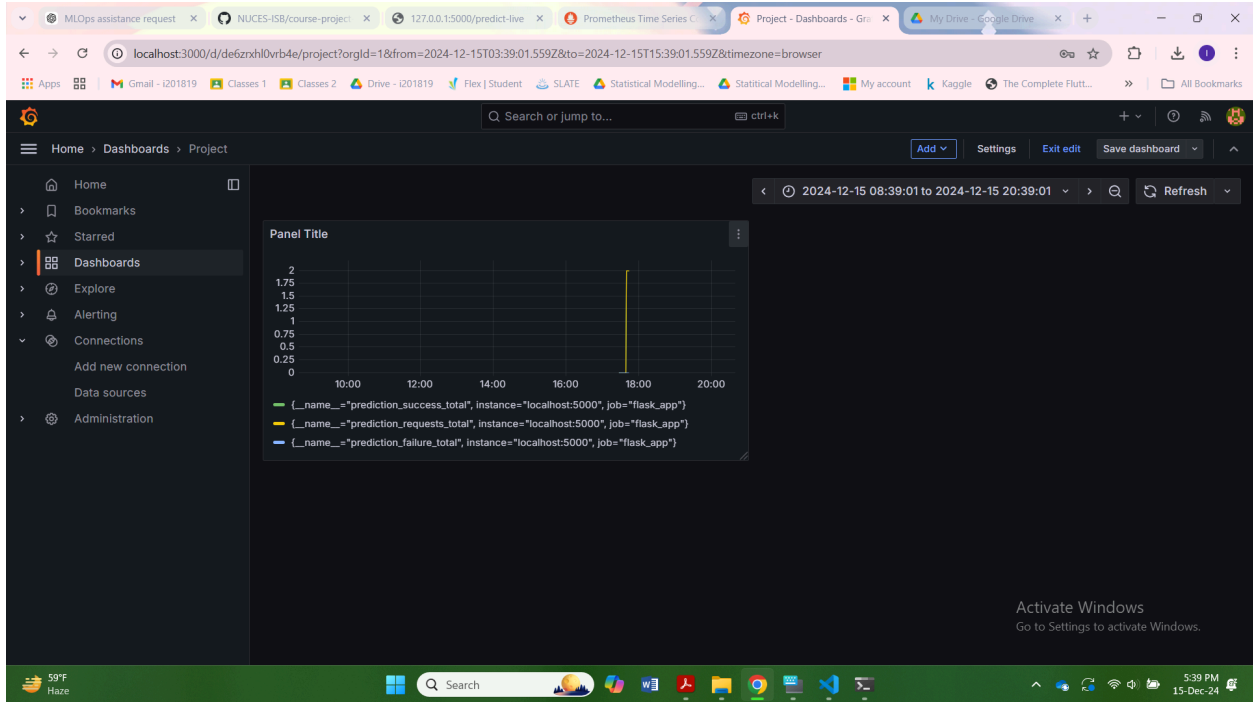
```
# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 4887.0
python_gc_objects_collected_total{generation="1"} 547.0
python_gc_objects_collected_total{generation="2"} 23.0
# HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 930.0
python_gc_collections_total{generation="1"} 84.0
python_gc_collections_total{generation="2"} 6.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="11",patchlevel="4",version="3.11.4"} 1.0
# HELP prediction_requests_total Total Prediction Requests
# TYPE prediction_requests_total counter
prediction_requests_total 0.0
# HELP prediction_requests_created Total Prediction Requests
# TYPE prediction_requests_created gauge
prediction_requests_created 1.7342649331371548e+09
# HELP prediction_success_total Successful Predictions
# TYPE prediction_success_total counter
prediction_success_total 0.0
# HELP prediction_success_created Successful Predictions
# TYPE prediction_success_created gauge
prediction_success_created 1.7342649331371548e+09
# HELP prediction_failure_total Failed Predictions
# TYPE prediction_failure_total counter
prediction_failure_total 0.0
# HELP prediction_failure_created Failed Predictions
# TYPE prediction_failure_created gauge
prediction_failure_created 1.7342649331371548e+09
```

Activate Windows  
Go to Settings to activate Windows.

USD/JPY  
+0.67%

Search

5:32 PM  
15-Dec-24



---

## Pipeline Optimization

Based on the monitored metrics:

- The API retry mechanism was enhanced to reduce failed prediction instances caused by connectivity issues.
- Prediction latency was minimized by optimizing model loading and input preprocessing.
- Prometheus and Grafana configurations were fine-tuned to improve the clarity of visualized metrics.

---

## Results and Insights

The integration of Prometheus and Grafana provided a comprehensive monitoring solution. Key results include:

- A fully functional real-time dashboard for tracking API performance and prediction metrics.
- Improved pipeline reliability and reduced failure rates.
- Enhanced understanding of system behavior under live conditions.

---

## Conclusion

Task 3 successfully validated the deployed pipeline under live data conditions. By leveraging Grafana and Prometheus, the system achieved robust monitoring and provided actionable insights for optimization. The resulting real-time dashboard serves as a critical tool for maintaining and improving pipeline performance in future deployments.

---